What goes into the model?

Both evaluations score the continuation after the cut. They give the model different amounts of prefix.

A · Length test: 128-position cap

256 documents · same first 32 continuation tokens at every cap

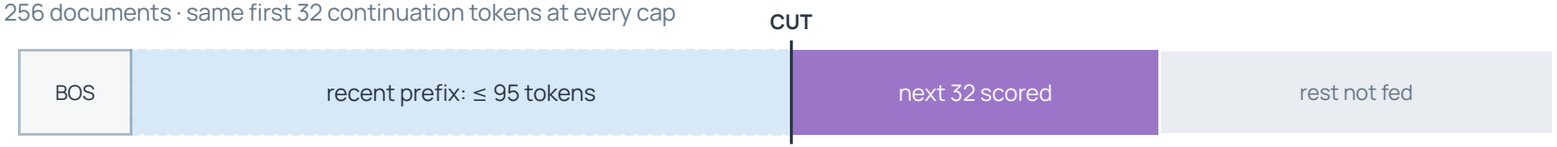

1 BOS + prefix + 32 scored tokens ≤ 128. Older prefix text is removed when needed.

5k → 20k mean gap: -0.0317 → -0.0321

B · Length test: 8,192-position cap

Same 256 documents and same 32 scored continuation tokens as A

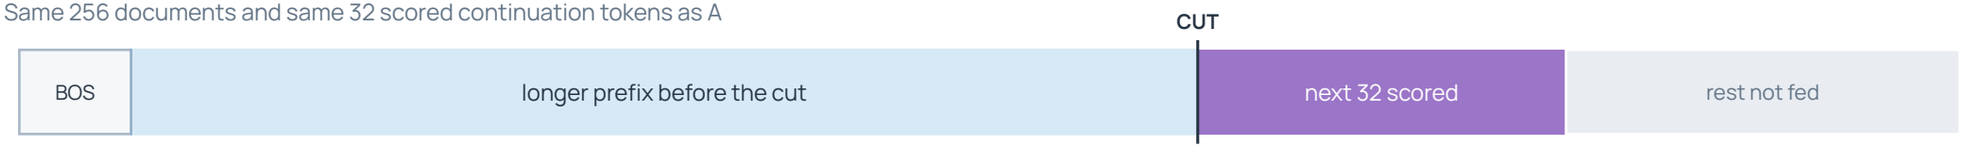

More earlier prefix text can fit. The cap is a maximum; actual inputs can be shorter.

5k → 20k mean gap: -0.0024 → +0.0234

C · Top right of zett_20k_transfer.pdf

1,936 documents · full saved prefix · first up to 128 continuation tokens scored

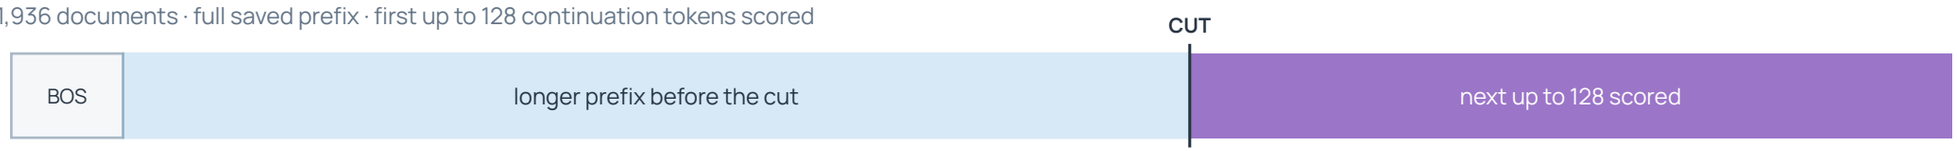

The full prefix plus BOS and scored continuation fits within 8,192 positions.

5k → 20k mean gap: +0.0280 → +0.0734

BLUE = prefix context (not scored) PURPLE = continuation (scored) GREY = not fed to the model · bars are schematic

A vs B changes prefix context while keeping the scored tokens fixed. C also changes the document set and scoring window.