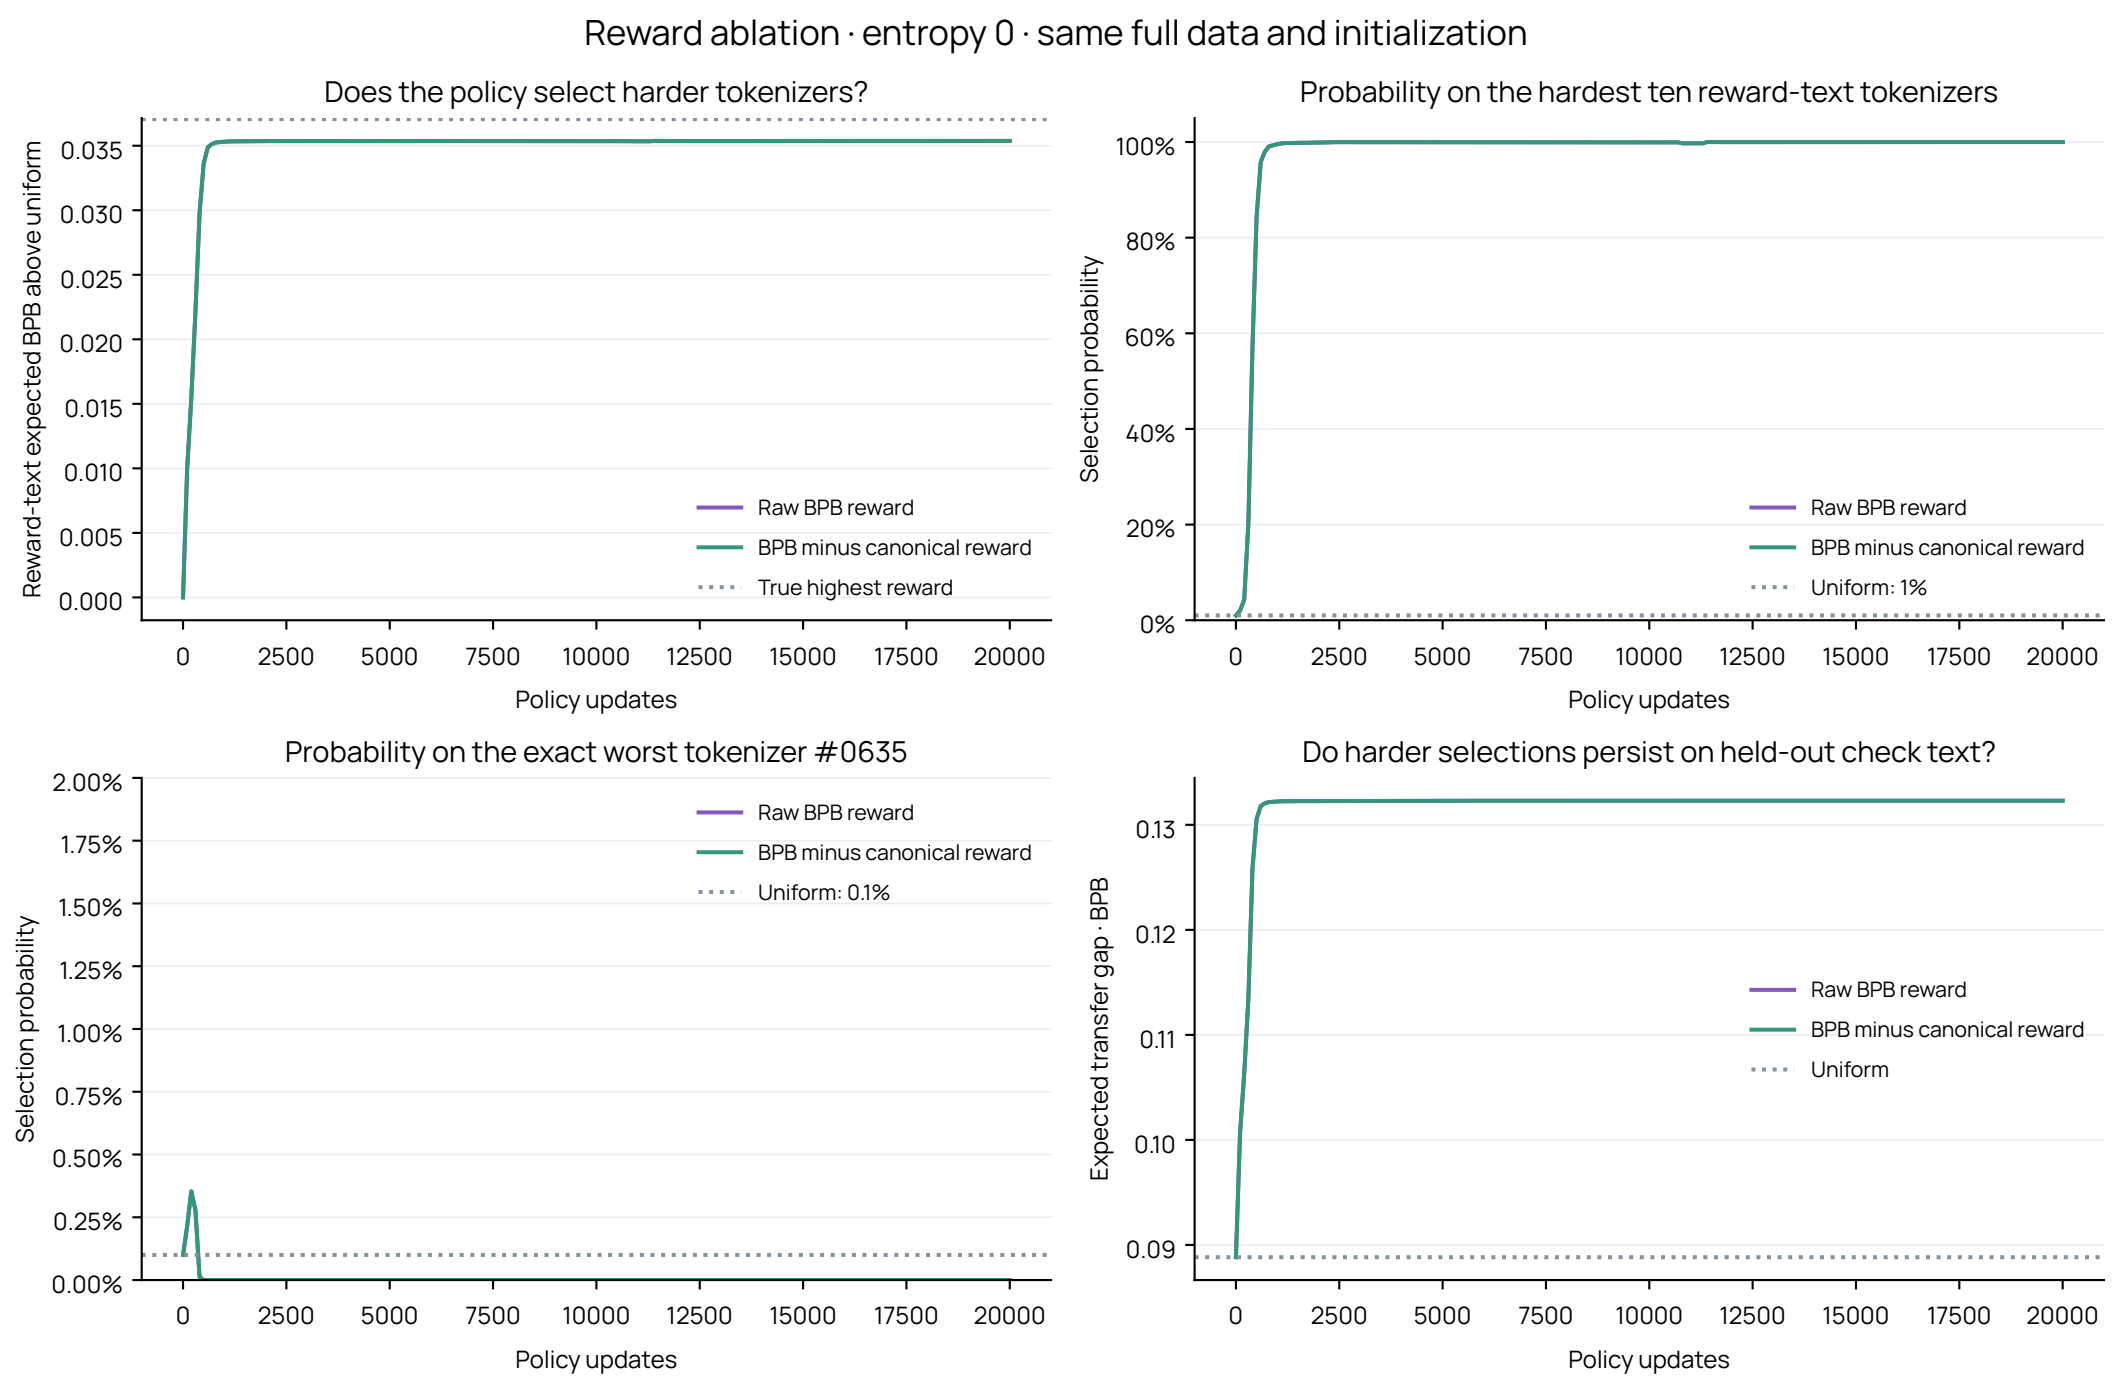
Reward ablation · entropy 0 · same full data and initialization
Does the policy select harder tokenizers?
Reward-text expected BPB above uniform
0.035
0.030
0.025
0.020
0.015
0.010
0.005
0.000
Raw BPB reward
BPB minus canonical reward
True highest reward
0
2500
5000
7500
10000
12500
15000
17500
20000
Policy updates
Probability on the hardest ten reward-text tokenizers
Selection probability
100%
80%
60%
40%
20%
0%
Raw BPB reward
BPB minus canonical reward
Uniform: 1%
Policy updates
Probability on the exact worst tokenizer #0635
Selection probability
2.00%
1.75%
1.50%
1.25%
1.00%
0.75%
0.50%
0.25%
0.00%
Raw BPB reward
BPB minus canonical reward
Uniform: 0.1%
Policy updates
Do harder selections persist on held-out check text?
Expected transfer gap · BPB
0.13
0.12
0.11
0.10
0.09
Raw BPB reward
BPB minus canonical reward
Uniform
Policy updates