RL20k · exhaustive training-tokenizer pool evaluation

1,000 native vocabularies · identical held-out text · 9 domains · next 32 canonical tokens · ≤128 positions

All 1,000 training-pool tokenizers

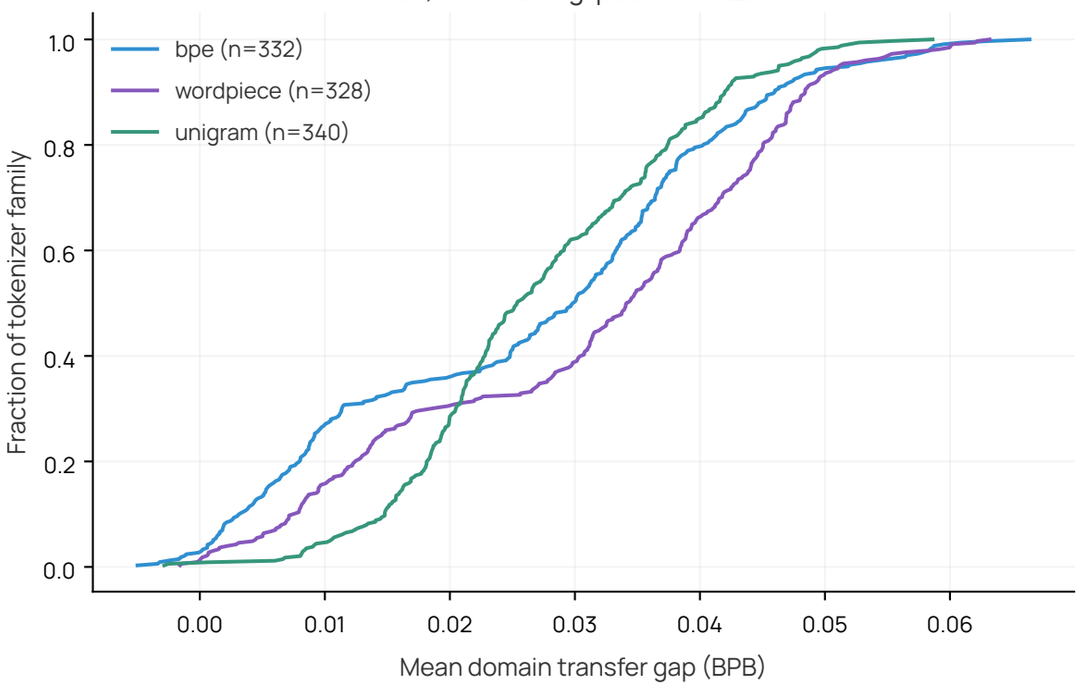

Ten worst screen candidates

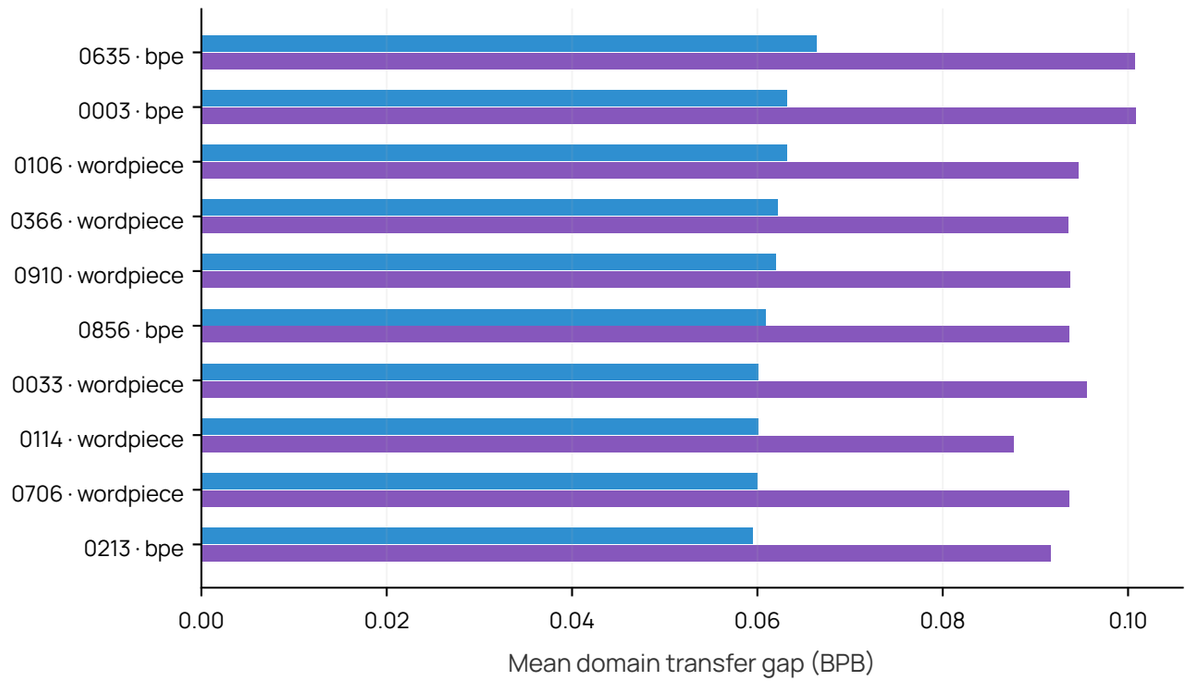

Where the worst candidates struggle

Training exposure and final transfer

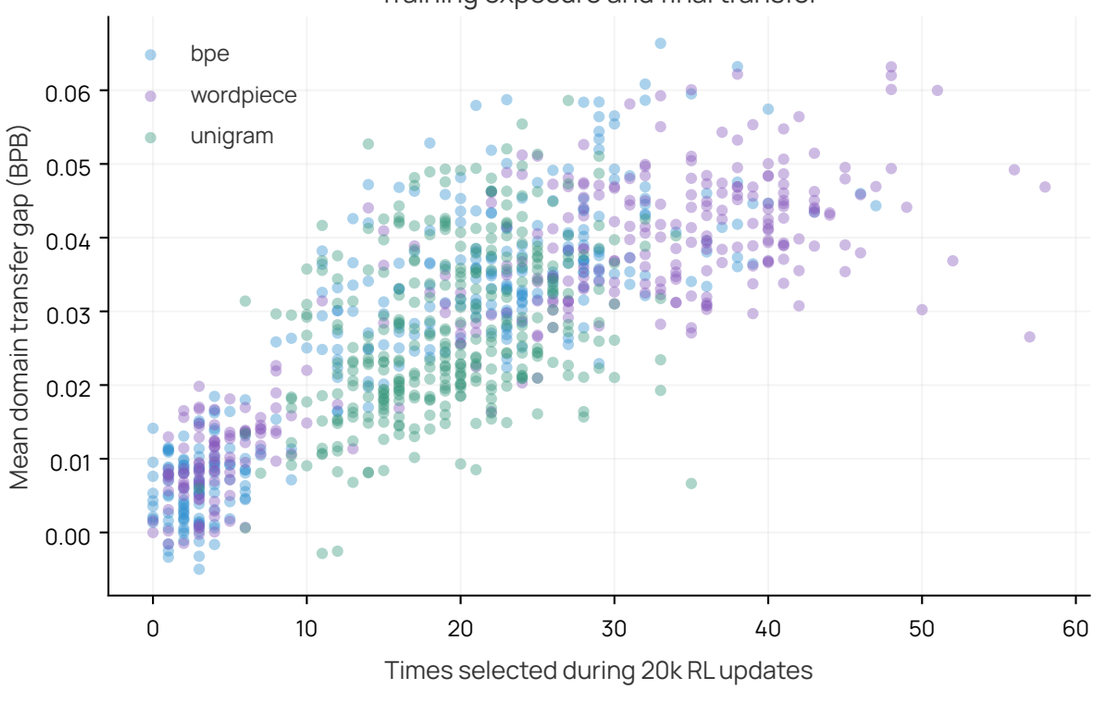

Higher gap is worse. Domains have equal weight; BPB within each domain is weighted by scored bytes.
Candidates selected on screen passages are checked on separate documents. One model-training seed; exposure correlations are descriptive.