Tokenizer coverage and reward audit · evidence overview

Existing 20k runs · no new training · 432 paired tokenizer/domain evaluations

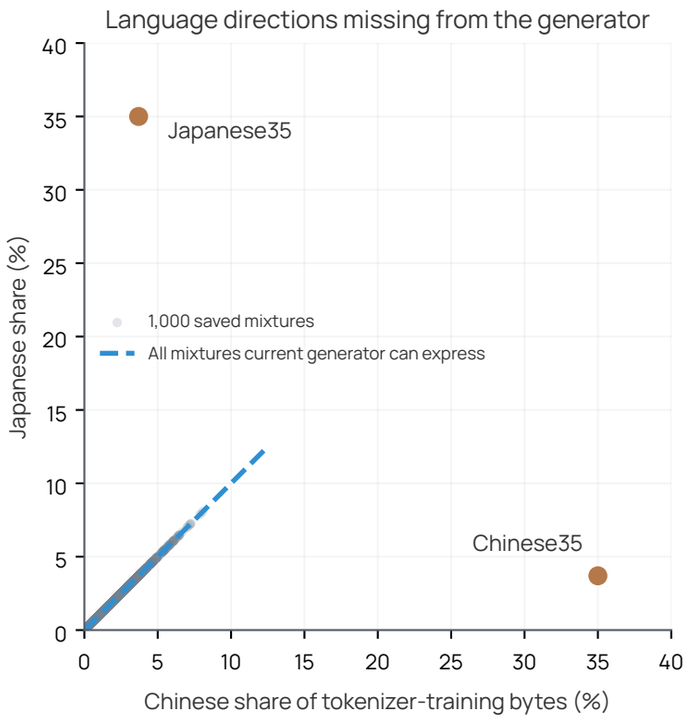

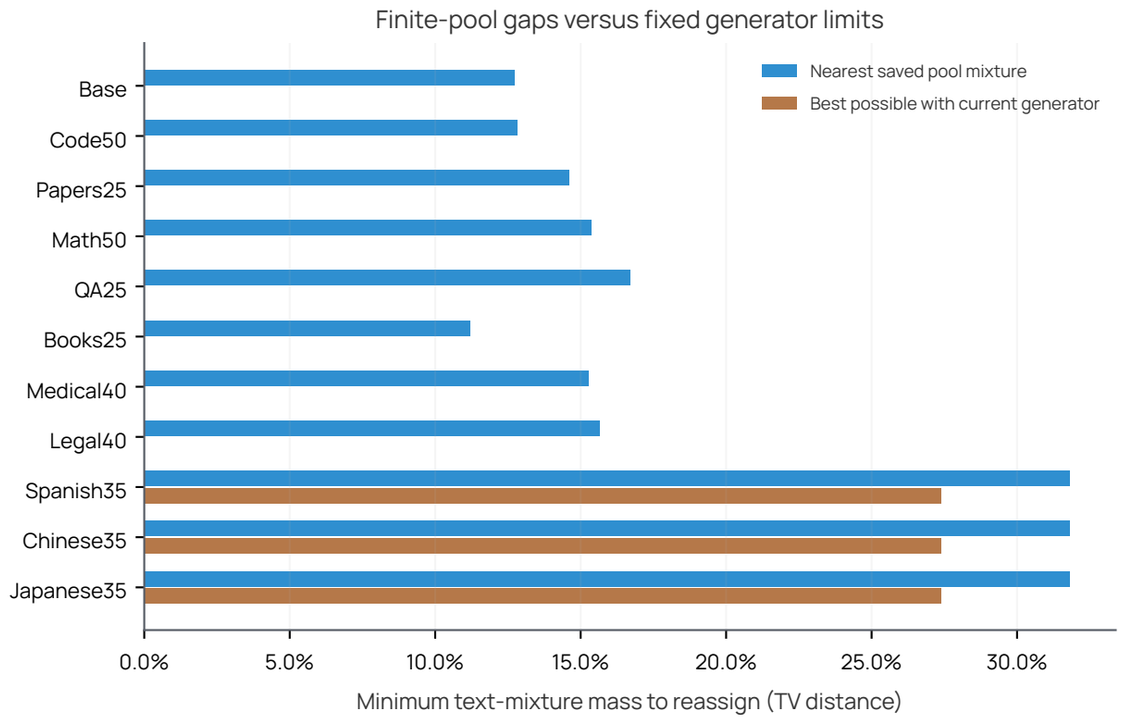

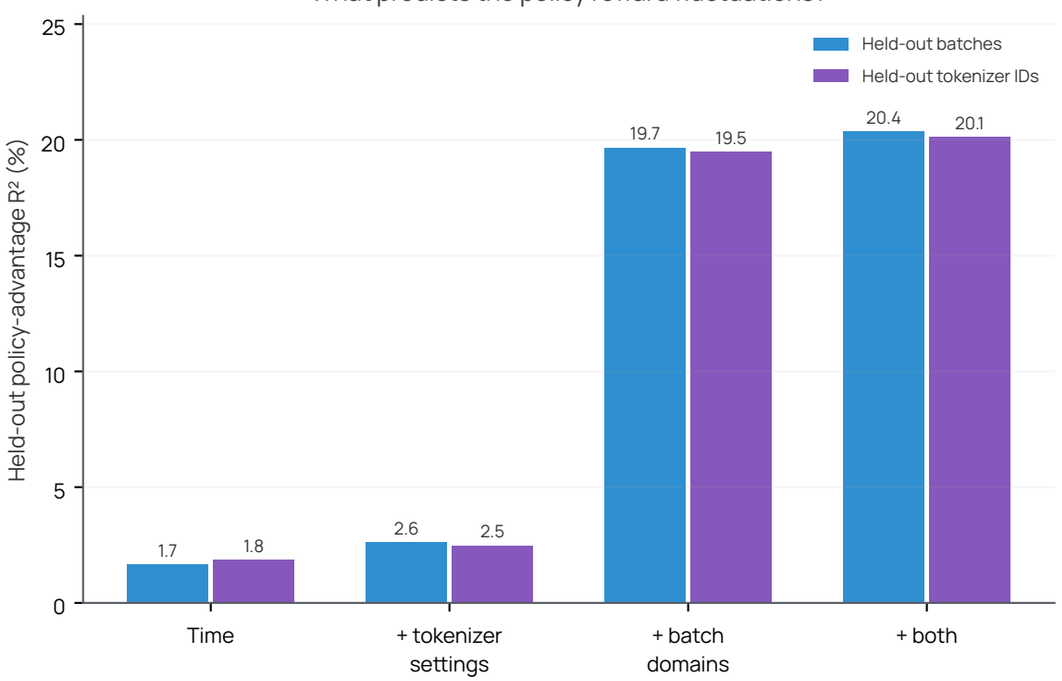

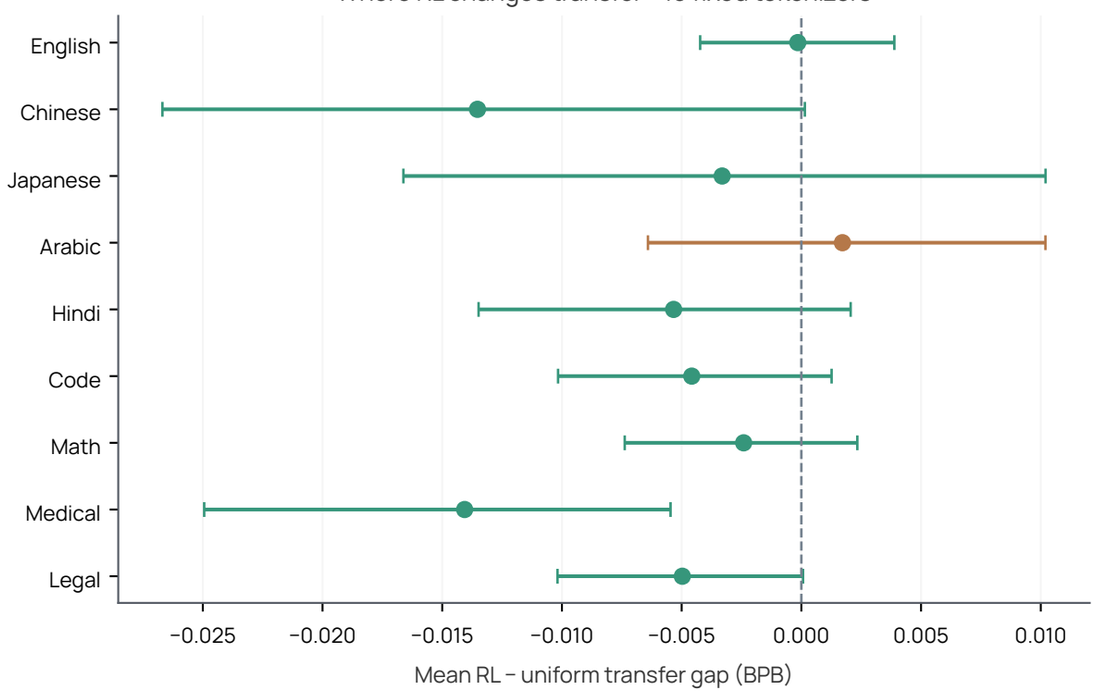

Mixture distances describe generator inputs, not a generalization bound. Reward prediction is descriptive.
Domain intervals resample shared source articles; tokenizers are fixed and utility-selected. One training seed per method; exploratory comparisons.