What the shared settings policy actually changed

Dotted lines show initial uniform group mass. Final policy and counts are read from cloud checkpoints.

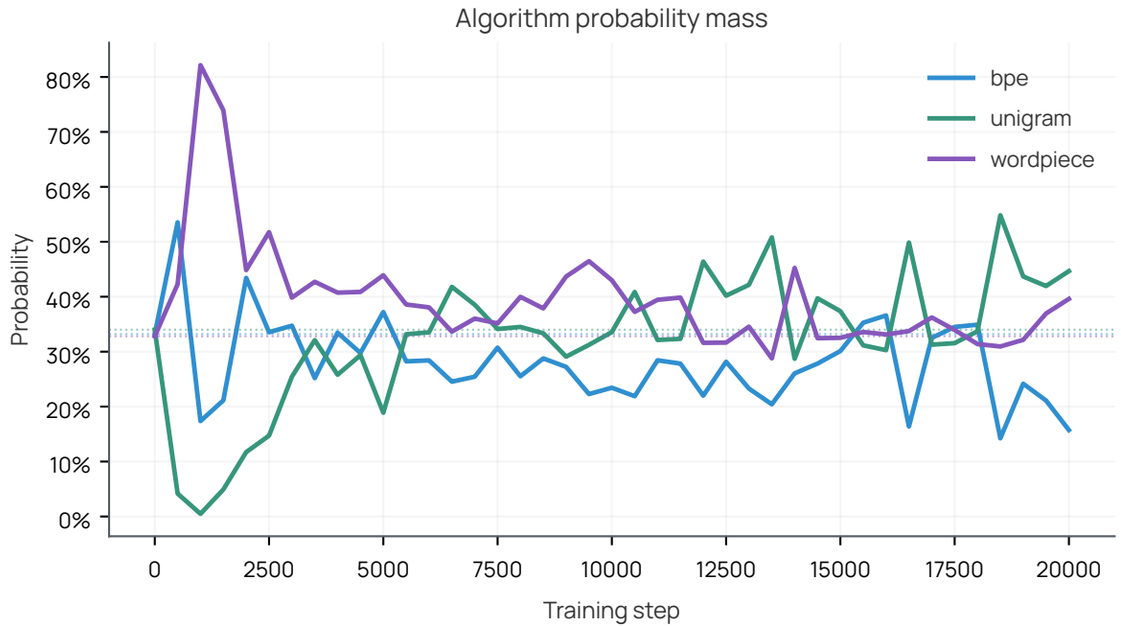

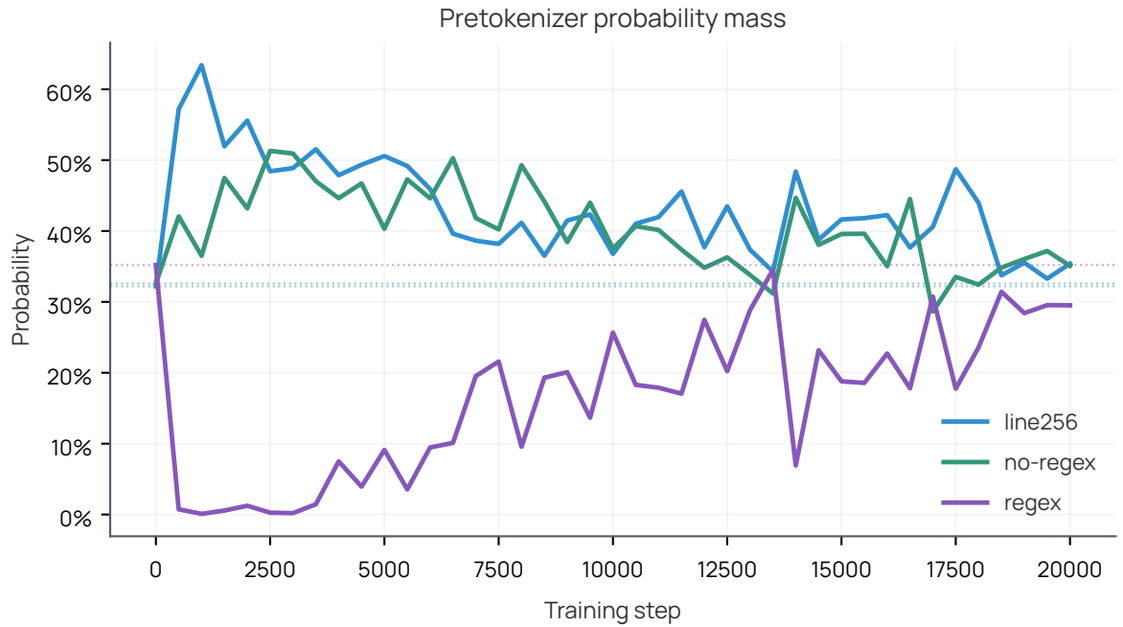

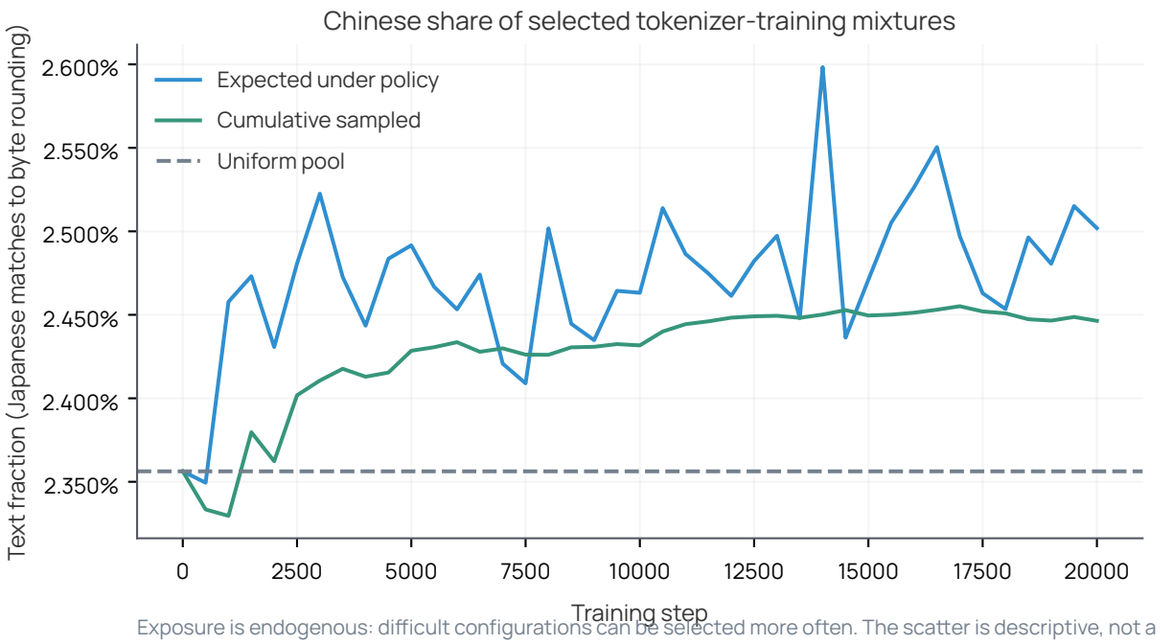

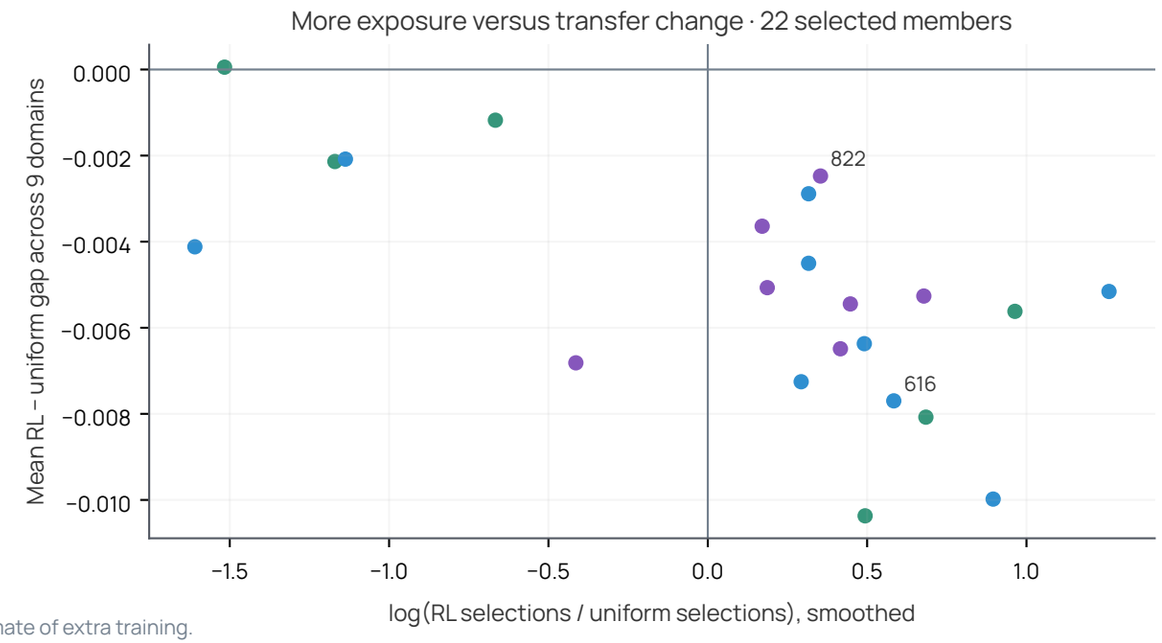

Exposure is endogenous: difficult configurations can be selected more often. The scatter is descriptive, not a causal estimate of extra training.
The Chinese/Japanese text fractions describe tokenizer-training corpora, not the LM's language mix.