RL and uniform · fixed sanity evaluations during training

Two fixed tokenizers · 92 examples each · ≤128 total positions · variable continuation length

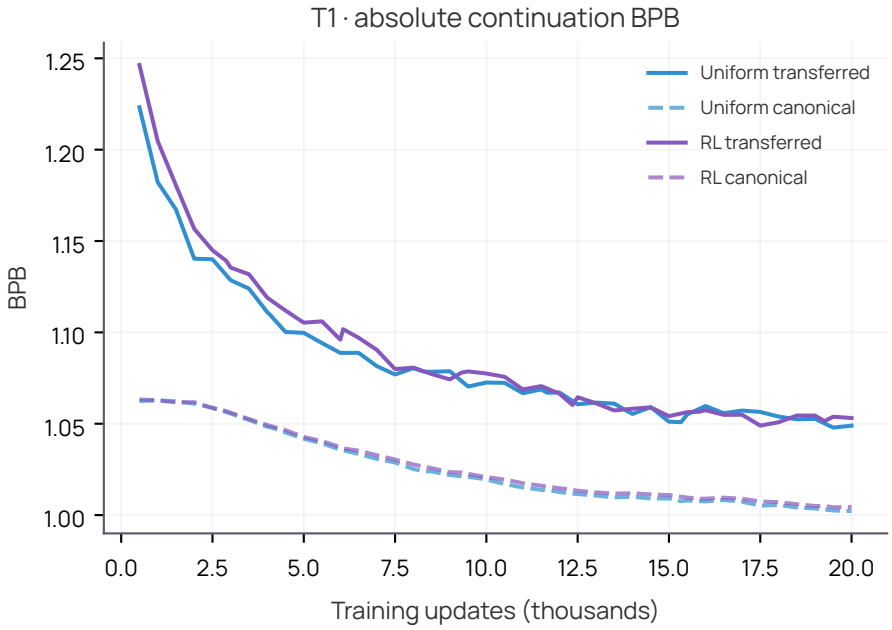

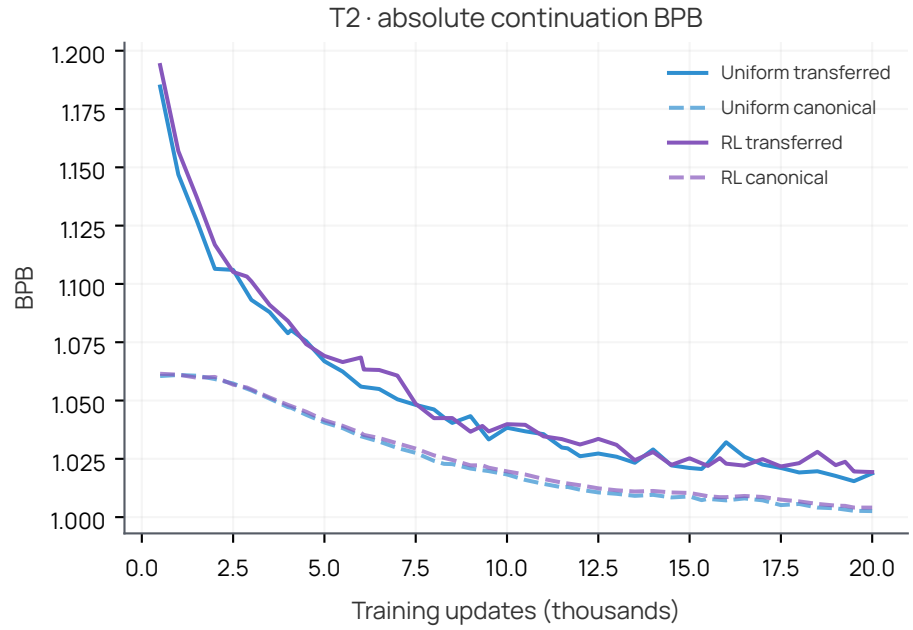

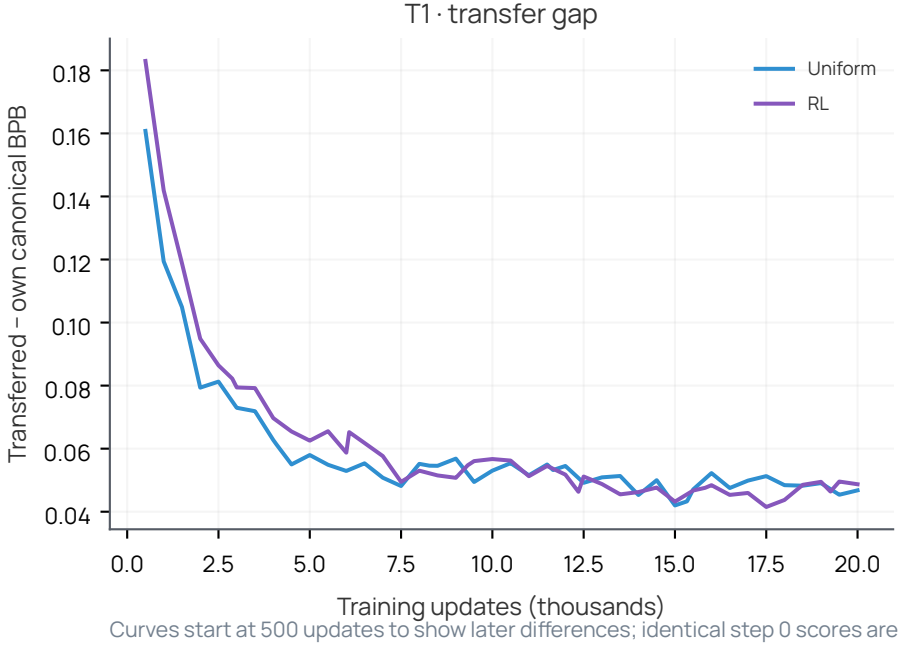

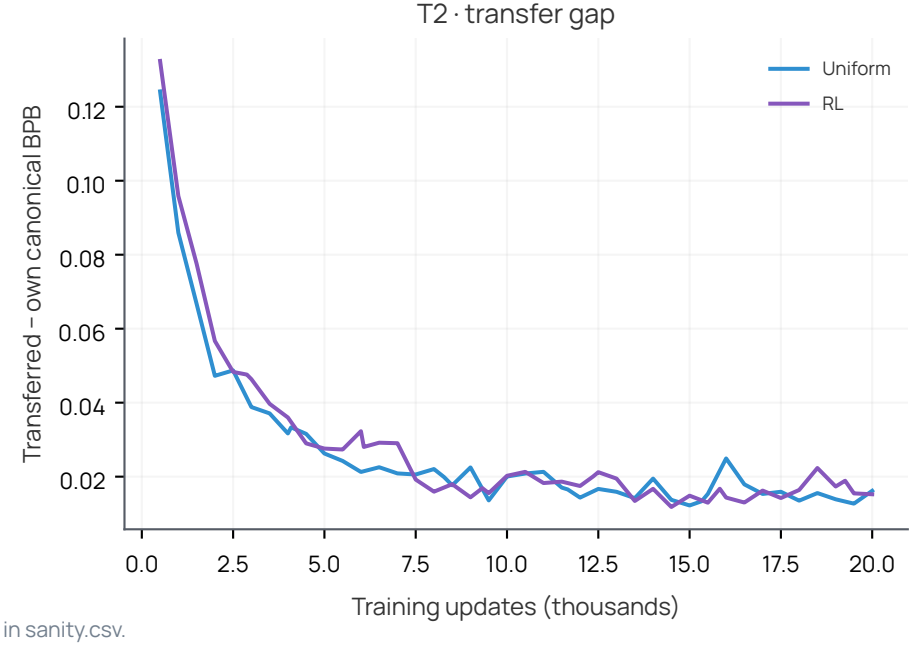

Curves start at 500 updates to show later differences; identical step 0 scores are included in sanity.csv.

These sanity curves are separate from the final 32-token OOD evaluations. Neither run establishes a consistent sanity-score advantage.